

# Assignment 2 - authored by, KARTHIKEYAN R

1. Download the dataset from the source [here \(https://drive.google.com/file/d/1\\_HcM0K8wt4b7Fusp=sharing\)](https://drive.google.com/file/d/1_HcM0K8wt4b7Fusp=sharing).

## About the dataset:

This dataset is all about churn modelling of a credit company. It has the details about the end user who are using credit card and also it has some variables to depict the churn of the customer.

**RowNumber** - Serial number of the rows

**CustomerId** - Unique identification of customer

**Surname** - Name of the customer

**CreditScore** - Cibil score of the customer

**Geography** - Location of the bank

**Gender** - Sex of the customer

**Age** - Age of the customer

**Tenure** - Repayment period for the credit amount

**Balance** - Current balance in their credit card

**NumOfProducts** - Products owned by the customer from the company

**HasCrCard** - Has credit card or not (0 - no, 1 - yes)

**IsActiveMember** - Is an active member or not

**EstimatedSalary** - Salary of the customer

**Exited** - Churn of the customer

In [1]:

```
import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

## 2. Load the dataset

In [3]:

```
df = pd.read_csv("Churn_Modelling.csv")
df.head()
```

Out[3]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	T
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

In [4]:

```
df.tail()
```

Out[4]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	T
9995	9996	15606229	Obijiaku	771	France	Male	38	
9996	9997	15569892	Johnstone	516	France	Male	38	
9997	9998	15584532	Liu	709	France	Female	38	
9998	9999	15682355	Sabbatini	772	Germany	Male	40	
9999	10000	15628319	Walker	792	France	Female	28	

### 3 a). Univariate analysis

In [5]:

```
#checking for categorical variables
category = df.select_dtypes(include=[np.object])
print("Categorical Variables: ",category.shape[1])

#checking for numerical variables
numerical = df.select_dtypes(include=[np.int64,np.float64])
print("Numerical Variables: ",numerical.shape[1])
```

Categorical Variables: 3  
Numerical Variables: 11

In [6]:

```
df.columns
```

Out[6]:

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

In [7]:

```
df.shape
```

Out[7]:


```
(10000, 14)
```

In [8]:

```
credit = df['CreditScore']  
credit.plot(kind="hist",figsize=(8,5))
```

Out[8]:

<AxesSubplot:ylabel='Frequency'>




In [9]:

```
geo = df['Geography'].value_counts()  
geo.plot(kind="pie",figsize=(10,8))
```

Out[9]:

<AxesSubplot:ylabel='Geography'>




In [10]:

```
sns.countplot(df['Gender'])
```

Out[10]:

<AxesSubplot:xlabel='Gender', ylabel='count'>




In [11]:

```
sns.distplot(df['Age'],hist=False)
```

Out[11]:

<AxesSubplot:xlabel='Age', ylabel='Density'>




In [12]:

```
plt.figure(figsize=(10,8))  
sns.countplot(df['Tenure'])
```

Out[12]:

<AxesSubplot:xlabel='Tenure', ylabel='count'>






In [13]:

```
product = df['NumOfProducts'].value_counts()  
product.plot(kind="pie",figsize=(10,8))
```

Out[13]:

<AxesSubplot:ylabel='NumOfProducts'>




In [14]:

```
cr = df['HasCrCard'].value_counts()  
cr.plot(kind="pie", figsize=(10,8))
```

Out[14]:

<AxesSubplot:ylabel='HasCrCard'>




In [15]:

```
plt.figure(figsize=(8,5))  
sns.countplot(df['IsActiveMember'])
```

Out[15]:

<AxesSubplot:xlabel='IsActiveMember', ylabel='count'>




In [16]:

```
sns.distplot(df['EstimatedSalary'], hist=False)
```

Out[16]:

<AxesSubplot: xlabel='EstimatedSalary', ylabel='Density'>




In [17]:

```
plt.figure(figsize=(8,5))  
sns.countplot(df['Exited'])
```

Out[17]:

<AxesSubplot:xlabel='Exited', ylabel='count'>



## Inference:

1. The data has 11 numerical variables and 3 categorical variables.
2. It has 10000 rows and 14 columns

3. The normalized credit score is around 700, More than 500 people have credit score greater than 800.
4. France occupies 50% of customers, where as Germany and Spain shared equal.
5. Dataset is dominated by Male Customers.
6. Median age is around 40 to 45.
7. Highest number of customer has thier tenure period for 2 years.
8. Credit company has maximum customers, who uses single product.
9. Most of the customer has credit card.
10. More than 40% of the population is not an active member.
11. The Churn is less compared to the satisfaction. **Dataset is imbalanced.**


## 3 b). Bivariate analysis

In [18]:

```
sns.barplot(x='Gender',y='CreditScore',hue='Geography',data=df)
```

Out[18]:

<AxesSubplot:xlabel='Gender', ylabel='CreditScore'>




In [19]:

```
sns.violinplot(x='Geography',y='Balance',data=df)
```

Out[19]:

<AxesSubplot:xlabel='Geography', ylabel='Balance'>




In [20]:

```
sns.violinplot(x='Geography',y='EstimatedSalary',data=df)
```

Out[20]:

<AxesSubplot:xlabel='Geography', ylabel='EstimatedSalary'>






In [21]:

```
sns.violinplot(x='Gender',y='Balance',data=df)
```

Out[21]:

<AxesSubplot:xlabel='Gender', ylabel='Balance'>




In [22]:

```
sns.barplot(x='Exited',y='CreditScore',hue='Gender',data=df)
```

Out[22]:

<AxesSubplot:xlabel='Exited', ylabel='CreditScore'>

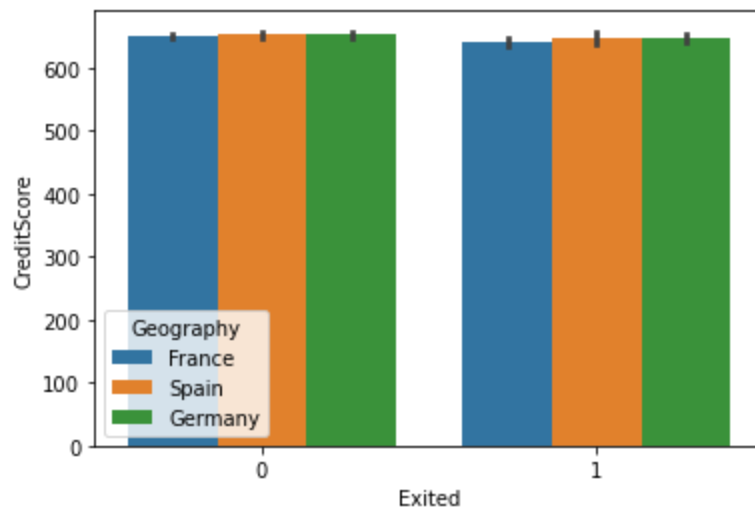


In [23]:

```
sns.barplot(x='Exited',y='CreditScore',hue='Geography',data=df)
```

Out[23]:

<AxesSubplot:xlabel='Exited', ylabel='CreditScore'>




In [24]:

```
sns.barplot(x='IsActiveMember',y='EstimatedSalary',hue='Gender',data=df)
```

Out[24]:

<AxesSubplot:xlabel='IsActiveMember', ylabel='EstimatedSalary'>




In [25]:

```
sns.barplot(x='Exited',y='Tenure',data=df)
```

Out[25]:

<AxesSubplot:xlabel='Exited', ylabel='Tenure'>



## Inference:

1. Credit score for Male is higher in Spain.
2. Average bank salary lies in the range of 100k to 150k.
3. Estimated salary is normalized and same for all country.
4. Credit score for churn is low.
5. Churn in Germany is higher compared to other countries.
6. Exited people tenure period is around 6 years.


## 3 c). Multivariate analysis

In [26]:

```
gp1 = df.groupby('Gender')['Geography'].value_counts()  
gp1.plot(kind='pie',figsize=(10,8))  
print(gp1)
```

Gender	Geography	
Female	France	2261
	Germany	1193
	Spain	1089
Male	France	2753
	Spain	1388
	Germany	1316

Name: Geography, dtype: int64



In [27]:

```
gp2 = df.groupby('Gender')['Age'].mean()  
print(gp2)
```

```
Gender  
Female    39.238389  
Male      38.658237  
Name: Age, dtype: float64
```

In [28]:

```
gp3 = df.groupby(['Gender', 'Geography'])['Tenure'].mean()  
print(gp3)
```


```
Gender  Geography  
Female  France      4.950022  
         Germany    4.965633  
         Spain      5.000000  
Male    France      5.049401  
         Germany    5.050152  
         Spain      5.057637  
Name: Tenure, dtype: float64
```

In [29]:

```
gp4 = df.groupby('Geography')['HasCrCard', 'IsActiveMember'].value_counts()
gp4.plot(kind="bar", figsize=(8,5))
print(gp4)
```

Geography	HasCrCard	IsActiveMember	
France	1	1	1826
		0	1717
	0	1	765
		0	706
Germany	1	0	918
		1	873
	0	1	375
		0	343
Spain	1	1	908
		0	813
	0	1	404
		0	352

dtype: int64






In [30]:

```
gp5 = df.groupby(['Gender', 'HasCrCard', 'IsActiveMember'])['EstimatedSalary'].me
gp5.plot(kind="line", figsize=(10,8))
print(gp5)
```

Gender	HasCrCard	IsActiveMember	
Female	0	0	102006.080352
		1	102648.996944
	1	0	101208.014567
		1	98510.152300
Male	0	0	99756.431151
		1	99873.931251
	1	0	100353.378996
		1	98914.378703

Name: EstimatedSalary, dtype: float64




In [31]:

```
gp6 = df.groupby(['Gender', 'IsActiveMember'])['Exited'].value_counts()  
gp6.plot(kind='bar', figsize=(10,8))  
print(gp6)
```

Gender	IsActiveMember	Exited	
Female	0	0	1534
		1	725
	1	0	1870
		1	414
Male	0	0	2013
		1	577
	1	0	2546
		1	321

Name: Exited, dtype: int64



In [32]:

```
gp7 = df.groupby('Exited')['Balance', 'EstimatedSalary'].mean()  
print(gp7)
```


	Balance	EstimatedSalary
Exited		
0	72745.296779	99738.391772
1	91108.539337	101465.677531

In [33]:

```
gp8 = df.groupby('Gender')['Geography','Exited'].value_counts()
gp8.plot(kind='bar',figsize=(10,8))
print (gp8)
```

Gender	Geography	Exited	
Female	France	0	1801
	Spain	0	858
	Germany	0	745
	France	1	460
	Germany	1	448
	Spain	1	231
Male	France	0	2403
	Spain	0	1206
	Germany	0	950
		1	366
	France	1	350
	Spain	1	182

dtype: int64



## Inference:

1. Germany has more female customers compared to male customers.
2. Average age of Male is 38, whereas average age of Female is 39.
3. Tenure period for both male and female is high in Spain.
4. It is observed that, those who have credit card are very active member in the company.
5. The estimated salary for a person who is not having credit card is high when compared to those having them.
6. Churn for inactive member is high compared to active member.
7. Those who churn has thier estimated salary very low.
8. France has the more churn rate.

## 4. Descriptive statistics

In [34]:

```
df.describe().T
```

Out[34]:

	count	mean	std	min	25%
<b>RowNumber</b>	10000.0	5.000500e+03	2886.895680	1.00	2500.75
<b>CustomerId</b>	10000.0	1.569094e+07	71936.186123	15565701.00	15628528.25
<b>CreditScore</b>	10000.0	6.505288e+02	96.653299	350.00	584.00
<b>Age</b>	10000.0	3.892180e+01	10.487806	18.00	32.00
<b>Tenure</b>	10000.0	5.012800e+00	2.892174	0.00	3.00
<b>Balance</b>	10000.0	7.648589e+04	62397.405202	0.00	0.00
<b>NumOfProducts</b>	10000.0	1.530200e+00	0.581654	1.00	1.00
<b>HasCrCard</b>	10000.0	7.055000e-01	0.455840	0.00	0.00
<b>IsActiveMember</b>	10000.0	5.151000e-01	0.499797	0.00	0.00
<b>EstimatedSalary</b>	10000.0	1.000902e+05	57510.492818	11.58	51002.11
<b>Exited</b>	10000.0	2.037000e-01	0.402769	0.00	0.00



## 5. Handling the missing values

In [35]:

```
df.isnull().sum()
```

Out[35]:

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited         0
dtype: int64
```

**There is no missing value in the dataset**


## 6. Finding outliers

In [36]:


```
def replace_outliers(df, field_name):
    Q1 = np.percentile(df[field_name],25,interpolation='midpoint')
    Q3 = np.percentile(df[field_name],75,interpolation='midpoint')
    IQR = Q3-Q1
    maxi = Q3+1.5*IQR
    mini = Q1-1.5*IQR
    df[field_name]=df[field_name].mask(df[field_name]>maxi,maxi)
    df[field_name]=df[field_name].mask(df[field_name]<mini,mini)
```

In [37]:

```
plt.title("Before removing outliers")
sns.boxplot(df['CreditScore'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'CreditScore')
sns.boxplot(df['CreditScore'])
plt.show()
```









In [38]:

```
plt.title("Before removing outliers")
sns.boxplot(df['Age'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'Age')
sns.boxplot(df['Age'])
plt.show()
```




In [39]:

```
sns.boxplot(df['Balance'])
```

Out[39]:

<AxesSubplot:xlabel='Balance'>




In [40]:

```
sns.boxplot(df['EstimatedSalary'])
```

Out[40]:

<AxesSubplot:xlabel='EstimatedSalary'>



Outliers from Age and Credit Score columns are removed

## 7. Check for categorical column and perform encoding.

In [41]:

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

In [42]:

```
df['Gender'] = le.fit_transform(df['Gender'])  
df['Geography'] = le.fit_transform(df['Geography'])
```

In [43]:

```
df.head()
```

Out[43]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619.0	0	0	42.0	2
1	2	15647311	Hill	608.0	2	0	41.0	1
2	3	15619304	Onio	502.0	0	0	42.0	8
3	4	15701354	Boni	699.0	0	0	39.0	1
4	5	15737888	Mitchell	850.0	2	0	43.0	2

Only two columns(Gender and Geography) is label encoded

## Removing unwanted columns and checking for feature importance

In [44]:

```
df = df.drop(['RowNumber', 'CustomerId', 'Surname'],axis=1)
```

In [45]:

```
df.head()
```

Out[45]:


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	IsActive
0	619.0	0	0	42.0	2	0.00	1	1
1	608.0	2	0	41.0	1	83807.86	1	1
2	502.0	0	0	42.0	8	159660.80	3	1
3	699.0	0	0	39.0	1	0.00	2	1
4	850.0	2	0	43.0	2	125510.82	1	1

In [46]:

```
plt.figure(figsize=(20,10))
df_lt = df.corr(method = "pearson")
df_lt1 = df_lt.where(np.tril(np.ones(df_lt.shape)).astype(np.bool))
sns.heatmap(df_lt1,annot=True,cmap="coolwarm")
```

Out[46]:

&lt;AxesSubplot:&gt;



1. The Removed columns are nothing to do with model building.
2. Feature importance also checked using pearson correlation.

## 8. Data Splitting

In [47]:

```
target = df['Exited']
data = df.drop(['Exited'],axis=1)
```

In [48]:

```
print(data.shape)
print(target.shape)
```

```
(10000, 10)
(10000,)
```

## 9. Scaling the independent values

In [49]:

```
from sklearn.preprocessing import StandardScaler
se = StandardScaler()
```

In [50]:

```
data['CreditScore'] = se.fit_transform(pd.DataFrame(data['CreditScore']))
data['Age'] = se.fit_transform(pd.DataFrame(data['Age']))
data['Balance'] = se.fit_transform(pd.DataFrame(data['Balance']))
data['EstimatedSalary'] = se.fit_transform(pd.DataFrame(data['EstimatedSalary']))
```

In [51]:

```
data.head()
```

Out[51]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProduct
0	-0.326878	0	0	0.342615	2	-1.225848	
1	-0.440804	2	0	0.240011	1	0.117350	
2	-1.538636	0	0	0.342615	8	1.333053	
3	0.501675	0	0	0.034803	1	-1.225848	
4	2.065569	2	0	0.445219	2	0.785728	



## 10. Train test split

In [52]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(data,target,test_size=0.25,ran
```

In [53]:

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(7500, 10)

(2500, 10)

(7500,)

(2500,)

## Conclusion:

1. The model is scaled using StandarScaler method.
2. The train and test split ratio is 15:5.
3. As it is a classification problem, basic algorithms can be used to build ML models.