

Assignment 2 - authored by, Kishore Akash YS

1. Download the dataset from the source here.

About the dataset:

This dataset is all about churn modelling of a credit company. It has the details about the end user who are using credit card and also it has some variables to depict the churn of the customer.

RowNumber - Serial number of the rows **CustomerId** - Unique identification of customer
Surname - Name of the customer **CreditScore** - Cibil score of the customer **Geography** - Location of the bank **Gender** - Sex of the customer **Age** - Age of the customer **Tenure** - Repayment period for the credit amount **Balance** - Current balance in their credit card
NumOfProducts - Products owned by the customer from the company **HasCrCard** - Has credit card or not (0 - no, 1 - yes) **IsActiveMember** - Is a active member or not
EstimatedSalary - Salary of the customer **Exited** - Churn of the customer

```
import warnings
warnings.filterwarnings("ignore")
```

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

2. Load the dataset

```
df = pd.read_csv("Churn_Modelling.csv")
df.head()
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age |
|---|-----------|------------|----------|-------------|-----------|--------|-----|
| 0 | 1 | 15634602 | Hargrave | 619 | France | Female | 42 |
| 1 | 2 | 15647311 | Hill | 608 | Spain | Female | 41 |
| 2 | 3 | 15619304 | Onio | 502 | France | Female | 42 |
| 3 | 4 | 15701354 | Boni | 699 | France | Female | 39 |
| 4 | 5 | 15737888 | Mitchell | 850 | Spain | Female | 43 |

| | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | |
|---|--------|---------|---------------|-----------|----------------|--|
| 0 | 2 | 0.00 | 1 | 1 | 1 | |

| | | | | | |
|---|---|-----------|---|---|---|
| 1 | 1 | 83807.86 | 1 | 0 | 1 |
| 2 | 8 | 159660.80 | 3 | 1 | 0 |
| 3 | 1 | 0.00 | 2 | 0 | 0 |
| 4 | 2 | 125510.82 | 1 | 1 | 1 |

| | EstimatedSalary | Exited |
|---|-----------------|--------|
| 0 | 101348.88 | 1 |
| 1 | 112542.58 | 0 |
| 2 | 113931.57 | 1 |
| 3 | 93826.63 | 0 |
| 4 | 79084.10 | 0 |

df.tail()

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender |
|-------|-----------|------------|-----------|-------------|-----------|--------|
| Age \ | | | | | | |
| 9995 | 9996 | 15606229 | Obijiaku | 771 | France | Male |
| 39 | | | | | | |
| 9996 | 9997 | 15569892 | Johnstone | 516 | France | Male |
| 35 | | | | | | |
| 9997 | 9998 | 15584532 | Liu | 709 | France | Female |
| 36 | | | | | | |
| 9998 | 9999 | 15682355 | Sabbatini | 772 | Germany | Male |
| 42 | | | | | | |
| 9999 | 10000 | 15628319 | Walker | 792 | France | Female |
| 28 | | | | | | |

| | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | \ |
|------|--------|-----------|---------------|-----------|----------------|---|
| 9995 | 5 | 0.00 | 2 | 1 | | 0 |
| 9996 | 10 | 57369.61 | 1 | 1 | | 1 |
| 9997 | 7 | 0.00 | 1 | 0 | | 1 |
| 9998 | 3 | 75075.31 | 2 | 1 | | 0 |
| 9999 | 4 | 130142.79 | 1 | 1 | | 0 |

| | EstimatedSalary | Exited |
|------|-----------------|--------|
| 9995 | 96270.64 | 0 |
| 9996 | 101699.77 | 0 |
| 9997 | 42085.58 | 1 |
| 9998 | 92888.52 | 1 |
| 9999 | 38190.78 | 0 |

3 a). Univariate analysis

#checking for categorical variables

```
category = df.select_dtypes(include=[np.object])
print("Categorical Variables: ",category.shape[1])
```

#checking for numerical variables

```
numerical = df.select_dtypes(include=[np.int64,np.float64])
print("Numerical Variables: ",numerical.shape[1])
```

```
Categorical Variables: 3  
Numerical Variables: 11
```

```
df.columns
```


```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',  
      'Geography',  
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',  
      'HasCrCard',  
      'IsActiveMember', 'EstimatedSalary', 'Exited'],  
      dtype='object')
```

```
df.shape
```

```
(10000, 14)
```


```
credit = df['CreditScore']  
credit.plot(kind="hist",figsize=(8,5))
```

```
<AxesSubplot:ylabel='Frequency'>
```




```
geo = df['Geography'].value_counts()  
geo.plot(kind="pie",figsize=(10,8))
```

```
<AxesSubplot:ylabel='Geography'>
```




```
sns.countplot(df['Gender'])  
<AxesSubplot:xlabel='Gender', ylabel='count'>
```




```
sns.distplot(df['Age'],hist=False)
```

```
<AxesSubplot:xlabel='Age', ylabel='Density'>
```




```
plt.figure(figsize=(10,8))  
sns.countplot(df['Tenure'])
```

```
<AxesSubplot:xlabel='Tenure', ylabel='count'>
```




```
product = df['NumOfProducts'].value_counts()  
product.plot(kind="pie", figsize=(10, 8))
```


```
<AxesSubplot:ylabel='NumOfProducts'>
```




```
cr = df['HasCrCard'].value_counts()
cr.plot(kind="pie",figsize=(10,8))
<AxesSubplot:ylabel='HasCrCard'>
```




```
plt.figure(figsize=(8,5))
sns.countplot(df['IsActiveMember'])
<AxesSubplot:xlabel='IsActiveMember', ylabel='count'>
```

```
sns.distplot(df['EstimatedSalary'], hist=False)  
<AxesSubplot: xlabel='EstimatedSalary', ylabel='Density'>
```



```
plt.figure(figsize=(8,5))  
sns.countplot(df['Exited'])  
<AxesSubplot: xlabel='Exited', ylabel='count'>
```




Inference:

1. The data has 11 numerical variables and 3 categorical variables.
2. It has 10000 rows and 14 columns
3. The normalized credit score is around 700, More than 500 people have credit score greater than 800.
4. France occupies 50% of customers, where as Germany and Spain shared equal.
5. Dataset is dominated by Male Customers.
6. Median age is around 40 to 45.
7. Highest number of customer has thier tenure period for 2 years.
8. Credit company has maximum customers, who uses single product.
9. Most of the customer has credit card.
10. More than 40% of the population is not an active member.
11. The Churn is less compared to the satisfaction. **Dataset is imbalanced.**


3 b). Bivariate analysis

```
sns.barplot(x='Gender',y='CreditScore',hue='Geography',data=df)
```


```
<AxesSubplot:xlabel='Gender', ylabel='CreditScore'>
```




```
sns.violinplot(x='Geography',y='Balance',data=df)
<AxesSubplot:xlabel='Geography', ylabel='Balance'>
```




```
sns.violinplot(x='Geography',y='EstimatedSalary',data=df)
<AxesSubplot:xlabel='Geography', ylabel='EstimatedSalary'>
```




```
sns.violinplot(x='Gender',y='Balance',data=df)
<AxesSubplot:xlabel='Gender', ylabel='Balance'>
```




```
sns.barplot(x='Exited',y='CreditScore',hue='Gender',data=df)
<AxesSubplot:xlabel='Exited', ylabel='CreditScore'>
```



```
sns.barplot(x='Exited',y='CreditScore',hue='Geography',data=df)
<AxesSubplot:xlabel='Exited', ylabel='CreditScore'>
```




```
sns.barplot(x='IsActiveMember',y='EstimatedSalary',hue='Gender',data=d
f)
<AxesSubplot:xlabel='IsActiveMember', ylabel='EstimatedSalary'>
```



```
sns.barplot(x='Exited',y='Tenure',data=df)
```

```
<AxesSubplot:xlabel='Exited', ylabel='Tenure'>
```



Inference:

1. Credit score for Male is higher in Spain.


2. Average bank salary lies in the range of 100k to 150k.
3. Estimated salary is normalized and same for all country.
4. Credit score for churn is low.
5. Churn in Germany is higher compared to other countries.
6. Exited people tenure period is around 6 years.

3 c). Multivariate analysis

```
gp1 = df.groupby('Gender')['Geography'].value_counts()  
gp1.plot(kind='pie',figsize=(10,8))  
print(gp1)
```

| Gender | Geography | |
|--------|-----------|------|
| Female | France | 2261 |
| | Germany | 1193 |
| | Spain | 1089 |
| Male | France | 2753 |
| | Spain | 1388 |
| | Germany | 1316 |

Name: Geography, dtype: int64



```
gp2 = df.groupby('Gender')['Age'].mean()
print(gp2)
```

```
Gender
Female    39.238389
Male      38.658237
Name: Age, dtype: float64
```

```
gp3 = df.groupby(['Gender', 'Geography'])['Tenure'].mean()
print(gp3)
```


```
Gender  Geography
Female  France      4.950022
        Germany     4.965633
        Spain      5.000000
Male    France      5.049401
        Germany     5.050152
        Spain      5.057637
Name: Tenure, dtype: float64
```



```
gp4 = df.groupby('Geography')
['HasCrCard', 'IsActiveMember'].value_counts()
gp4.plot(kind="bar", figsize=(8,5))
print(gp4)
```

| Geography | HasCrCard | IsActiveMember | |
|-----------|-----------|----------------|------|
| France | 1 | 1 | 1826 |
| | | 0 | 1717 |
| | 0 | 1 | 765 |
| | | 0 | 706 |
| Germany | 1 | 0 | 918 |
| | | 1 | 873 |
| | 0 | 1 | 375 |
| | | 0 | 343 |
| Spain | 1 | 1 | 908 |
| | | 0 | 813 |
| | 0 | 1 | 404 |
| | | 0 | 352 |


dtype: int64



```
gp5 = df.groupby(['Gender', 'HasCrCard', 'IsActiveMember'])
['EstimatedSalary'].mean()
gp5.plot(kind="line", figsize=(10,8))
print(gp5)
```

| Gender | HasCrCard | IsActiveMember | |
|--------|-----------|----------------|---------------|
| Female | 0 | 0 | 102006.080352 |
| | | 1 | 102648.996944 |
| | 1 | 0 | 101208.014567 |
| | | 1 | 98510.152300 |
| Male | 0 | 0 | 99756.431151 |
| | | 1 | 99873.931251 |
| | 1 | 0 | 100353.378996 |
| | | 1 | 98914.378703 |


Name: EstimatedSalary, dtype: float64



```
gp6 = df.groupby(['Gender', 'IsActiveMember'])['Exited'].value_counts()
gp6.plot(kind='bar', figsize=(10,8))
print(gp6)
```

| Gender | IsActiveMember | Exited | |
|--------|----------------|--------|------|
| Female | 0 | 0 | 1534 |
| | | 1 | 725 |
| | 1 | 0 | 1870 |
| | | 1 | 414 |
| Male | 0 | 0 | 2013 |
| | | 1 | 577 |
| | 1 | 0 | 2546 |
| | | 1 | |

1 321
Name: Exited, dtype: int64



```
gp7 = df.groupby('Exited')['Balance', 'EstimatedSalary'].mean()
print(gp7)
```


| | Balance | EstimatedSalary |
|--------|--------------|-----------------|
| Exited | | |
| 0 | 72745.296779 | 99738.391772 |
| 1 | 91108.539337 | 101465.677531 |

```
gp8 = df.groupby('Gender')['Geography', 'Exited'].value_counts()
gp8.plot(kind='bar', figsize=(10,8))
print (gp8)
```

| Gender | Geography | Exited | |
|--------|-----------|--------|------|
| Female | France | 0 | 1801 |
| | Spain | 0 | 858 |
| | Germany | 0 | 745 |
| | France | 1 | 460 |
| | Germany | 1 | 448 |

| | | | |
|------|---------|---|------|
| | Spain | 1 | 231 |
| Male | France | 0 | 2403 |
| | Spain | 0 | 1206 |
| | Germany | 0 | 950 |
| | | 1 | 366 |
| | France | 1 | 350 |
| | Spain | 1 | 182 |

dtype: int64



Inference:

1. Germany has more female customers compared to male customers.
2. Average age of Male is 38, whereas average age of Female is 39.
3. Tenure period for both male and female is high in Spain.
4. It is observed that, those who have credit card are very active member in the company.

5. The estimated salary for a person who is not having credit card is high when compared to those having them.
6. Churn for inactive member is high compared to active member.
7. Those who churn has thier estimated salary very low.
8. France has the more churn rate.

4. Descriptive statistics

df.describe().T

| | count | mean | std | min \ |
|-----------------|---------|--------------|--------------|-------------|
| RowNumber | 10000.0 | 5.000500e+03 | 2886.895680 | 1.00 |
| CustomerId | 10000.0 | 1.569094e+07 | 71936.186123 | 15565701.00 |
| CreditScore | 10000.0 | 6.505288e+02 | 96.653299 | 350.00 |
| Age | 10000.0 | 3.892180e+01 | 10.487806 | 18.00 |
| Tenure | 10000.0 | 5.012800e+00 | 2.892174 | 0.00 |
| Balance | 10000.0 | 7.648589e+04 | 62397.405202 | 0.00 |
| NumOfProducts | 10000.0 | 1.530200e+00 | 0.581654 | 1.00 |
| HasCrCard | 10000.0 | 7.055000e-01 | 0.455840 | 0.00 |
| IsActiveMember | 10000.0 | 5.151000e-01 | 0.499797 | 0.00 |
| EstimatedSalary | 10000.0 | 1.000902e+05 | 57510.492818 | 11.58 |
| Exited | 10000.0 | 2.037000e-01 | 0.402769 | 0.00 |

| | 25% | 50% | 75% | max |
|-----------------|-------------|--------------|--------------|-------------|
| RowNumber | 2500.75 | 5.000500e+03 | 7.500250e+03 | 10000.00 |
| CustomerId | 15628528.25 | 1.569074e+07 | 1.575323e+07 | 15815690.00 |
| CreditScore | 584.00 | 6.520000e+02 | 7.180000e+02 | 850.00 |
| Age | 32.00 | 3.700000e+01 | 4.400000e+01 | 92.00 |
| Tenure | 3.00 | 5.000000e+00 | 7.000000e+00 | 10.00 |
| Balance | 0.00 | 9.719854e+04 | 1.276442e+05 | 250898.09 |
| NumOfProducts | 1.00 | 1.000000e+00 | 2.000000e+00 | 4.00 |
| HasCrCard | 0.00 | 1.000000e+00 | 1.000000e+00 | 1.00 |
| IsActiveMember | 0.00 | 1.000000e+00 | 1.000000e+00 | 1.00 |
| EstimatedSalary | 51002.11 | 1.001939e+05 | 1.493882e+05 | 199992.48 |
| Exited | 0.00 | 0.000000e+00 | 0.000000e+00 | 1.00 |

5. Handling the missing values

```
df.isnull().sum()
```



```
RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender         0
Age           0
Tenure         0
Balance        0
NumOfProducts 0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

There is no missing value in the dataset

6. Finding outliers



```
def replace_outliers(df, field_name):
    Q1 = np.percentile(df[field_name],25,interpolation='midpoint')
    Q3 = np.percentile(df[field_name],75,interpolation='midpoint')
    IQR = Q3-Q1
    maxi = Q3+1.5*IQR
    mini = Q1-1.5*IQR
    df[field_name]=df[field_name].mask(df[field_name]>maxi,maxi)
    df[field_name]=df[field_name].mask(df[field_name]<mini,mini)

plt.title("Before removing outliers")
sns.boxplot(df['CreditScore'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'CreditScore')
sns.boxplot(df['CreditScore'])
plt.show()
```




```
plt.title("Before removing outliers")
sns.boxplot(df['Age'])
plt.show()
plt.title("After removing outliers")
replace_outliers(df, 'Age')
```


```
sns.boxplot(df['Age'])  
plt.show()
```



```
sns.boxplot(df['Balance'])  
<AxesSubplot:xlabel='Balance'>
```

```
sns.boxplot(df['EstimatedSalary'])  
<AxesSubplot:xlabel='EstimatedSalary'>
```



Outliers from Age and Credit Score columns are removed

7. Check for categorical column and perform encoding.

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
df['Gender'] = le.fit_transform(df['Gender'])  
df['Geography'] = le.fit_transform(df['Geography'])
```

```
df.head()
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender |
|-------|-----------|------------|----------|-------------|-----------|--------|
| Age \ | | | | | | |
| 0 | 1 | 15634602 | Hargrave | 619.0 | 0 | 0 |
| 42.0 | | | | | | |
| 1 | 2 | 15647311 | Hill | 608.0 | 2 | 0 |
| 41.0 | | | | | | |
| 2 | 3 | 15619304 | Onio | 502.0 | 0 | 0 |
| 42.0 | | | | | | |
| 3 | 4 | 15701354 | Boni | 699.0 | 0 | 0 |
| 39.0 | | | | | | |
| 4 | 5 | 15737888 | Mitchell | 850.0 | 2 | 0 |
| 43.0 | | | | | | |

| | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | \ |
|---|--------|-----------|---------------|-----------|----------------|---|
| 0 | 2 | 0.00 | 1 | 1 | 1 | |
| 1 | 1 | 83807.86 | 1 | 0 | 1 | |
| 2 | 8 | 159660.80 | 3 | 1 | 0 | |
| 3 | 1 | 0.00 | 2 | 0 | 0 | |
| 4 | 2 | 125510.82 | 1 | 1 | 1 | |

| | EstimatedSalary | Exited |
|---|-----------------|--------|
| 0 | 101348.88 | 1 |
| 1 | 112542.58 | 0 |
| 2 | 113931.57 | 1 |
| 3 | 93826.63 | 0 |
| 4 | 79084.10 | 0 |

Only two columns(Gender and Geography) is label encoded

Removing unwanted columns and checking for feature importance

```
df = df.drop(['RowNumber', 'CustomerId', 'Surname'],axis=1)
```

```
df.head()
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance |
|-----------------|-------------|-----------|--------|------|--------|----------|
| NumOfProducts \ | | | | | | |
| 0 | 619.0 | 0 | 0 | 42.0 | 2 | 0.00 |
| 1 | | | | | | |
| 1 | 608.0 | 2 | 0 | 41.0 | 1 | 83807.86 |
| 1 | | | | | | |

```

2          502.0          0          0  42.0          8  159660.80
3
3          699.0          0          0  39.0          1          0.00
2
4          850.0          2          0  43.0          2  125510.82
1

```

```

      HasCrCard  IsActiveMember  EstimatedSalary  Exited
0             1             1         101348.88         1
1             0             1         112542.58         0
2             1             0         113931.57         1
3             0             0          93826.63         0
4             1             1          79084.10         0


```

```

plt.figure(figsize=(20,10))
df_lt = df.corr(method = "pearson")
df_lt1 = df_lt.where(np.tril(np.ones(df_lt.shape)).astype(np.bool))
sns.heatmap(df_lt1,annot=True,cmap="coolwarm")

```

<AxesSubplot:>



1. The Removed columns are nothing to do with model building. 2. Feature importance also checked using pearson correlation.

8. Data Splitting

```

target = df['Exited']
data = df.drop(['Exited'],axis=1)

print(data.shape)
print(target.shape)

```

```
(10000, 10)
(10000,)
```

9. Scaling the independent values

```
from sklearn.preprocessing import StandardScaler
se = StandardScaler()

data['CreditScore'] =
se.fit_transform(pd.DataFrame(data['CreditScore']))
data['Age'] = se.fit_transform(pd.DataFrame(data['Age']))
data['Balance'] = se.fit_transform(pd.DataFrame(data['Balance']))
data['EstimatedSalary'] =
se.fit_transform(pd.DataFrame(data['EstimatedSalary']))

data.head()
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance |
|---|-------------|-----------|--------|----------|--------|-----------|
| 0 | -0.326878 | 0 | 0 | 0.342615 | 2 | -1.225848 |
| 1 | | | | | | |
| 1 | -0.440804 | 2 | 0 | 0.240011 | 1 | 0.117350 |
| 1 | | | | | | |
| 2 | -1.538636 | 0 | 0 | 0.342615 | 8 | 1.333053 |
| 3 | | | | | | |
| 3 | 0.501675 | 0 | 0 | 0.034803 | 1 | -1.225848 |
| 2 | | | | | | |
| 4 | 2.065569 | 2 | 0 | 0.445219 | 2 | 0.785728 |
| 1 | | | | | | |

| | HasCrCard | IsActiveMember | EstimatedSalary |
|---|-----------|----------------|-----------------|
| 0 | 1 | 1 | 0.021886 |
| 1 | 0 | 1 | 0.216534 |
| 2 | 1 | 0 | 0.240687 |
| 3 | 0 | 0 | -0.108918 |
| 4 | 1 | 1 | -0.365276 |

10. Train test split

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(data,target,test_size=0.25,random_state=101)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(7500, 10)
(2500, 10)
```

(7500,)
(2500,)

Conclusion:

1. The model is scaled using StandarScaler method.
2. The train and test split ratio is 15:5.
3. As it is a classification problem, basic algorithms can be used to build ML models.