

Assignment - 3, authored by Kishore Akash YS

1. Download the dataset [here \(https://drive.google.com/file/d/1xkynpL15pt6KT3YSIDimu4A5iRU9qYck/view\)](https://drive.google.com/file/d/1xkynpL15pt6KT3YSIDimu4A5iRU9qYck/view).

Importing necessary Libraries

```
In [1]: import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, Dropout, Conv2D, Flatten, Ma
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.applications.resnet50 import preprocess_input
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, i
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
```

Data Augumentation

- Dataset consist of 5 classes.
- **Daisy** - European Species of Aster family.
- **Sunflower** - Identified as the genus of Helianthus.
- **Tulip** - Belong to the species of spring blooming geophytes.
- **Rose** - Belongs to the family of rosaceae.
- **Dandelion** - Indentifies as the genus of Asterceae.

```
In [3]: path = 'flowers/'
```

```
In [4]: train_data_gen = ImageDataGenerator(rescale = 1./255,
                                             shear_range = 0.2,
                                             zoom_range = 0.2,
                                             horizontal_flip = True,
                                             validation_split = 0.30)
test_data_gen = ImageDataGenerator(rescale = 1./255, validation_split = 0.30)
```

```
In [5]: training_set = train_data_gen.flow_from_directory(path,
                                                         target_size=(64,64),
                                                         batch_size=100,
                                                         class_mode='categorical',
                                                         shuffle=True,
                                                         color_mode='rgb',
                                                         subset = 'training')

testing_set = test_data_gen.flow_from_directory(path,
                                                target_size=(64,64),
                                                batch_size=100,
                                                class_mode='categorical',
                                                shuffle=True,
                                                color_mode='rgb',
                                                subset = 'validation')
```

Found 3024 images belonging to 5 classes.
Found 1293 images belonging to 5 classes.

Model building using CNN

1. Create the model

```
In [6]: model = Sequential()

#convolution and Pooling layer 1
model.add(Conv2D(filters=48,kernel_size=3,activation='relu',input_shape=(64,64
model.add(MaxPool2D(pool_size=2,strides=2))
model.add(Dropout(0.2))

#convolution and Pooling layer 2
model.add(Conv2D(filters=32,kernel_size=3,activation='relu'))
model.add(MaxPool2D(pool_size=2,strides=2))
model.add(Dropout(0.2))

#Flattening the images
model.add(Flatten())

#Fully Connected Layers
model.add(Dense(64,activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(5,activation='softmax'))
```

In [7]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 48)	1344
max_pooling2d (MaxPooling2D)	(None, 31, 31, 48)	0
dropout (Dropout)	(None, 31, 31, 48)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	13856
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
dropout_1 (Dropout)	(None, 14, 14, 32)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 64)	401472
dropout_2 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 5)	325
=====		
Total params: 416,997		
Trainable params: 416,997		
Non-trainable params: 0		

2. Compile the Model

In [8]: `model.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accu`

3. Adding callbacks to avoid overfitting

```
In [9]: early_stop = EarlyStopping(monitor='val_accuracy',
                                   patience=5,verbose=1,mode='auto')

lr = ReduceLROnPlateau(monitor='val_accuracy',
                       factor=0.2,patience=5,
                       min_lr=0.00001)

callback = [early_stop,lr]
```

4. Training the Model


```
In [10]: result = model.fit(x=training_set, validation_data=testing_set, epochs=50)
```

```
Epoch 1/50
31/31 [=====] - 17s 536ms/step - loss: 1.4674 - accuracy: 0.3532 - val_loss: 1.2853 - val_accuracy: 0.4200
Epoch 2/50
31/31 [=====] - 18s 598ms/step - loss: 1.2396 - accuracy: 0.4580 - val_loss: 1.1973 - val_accuracy: 0.4942
Epoch 3/50
31/31 [=====] - 17s 544ms/step - loss: 1.1910 - accuracy: 0.4825 - val_loss: 1.1977 - val_accuracy: 0.5004
Epoch 4/50
31/31 [=====] - 20s 628ms/step - loss: 1.1573 - accuracy: 0.5136 - val_loss: 1.1799 - val_accuracy: 0.5004
Epoch 5/50
31/31 [=====] - 17s 556ms/step - loss: 1.1126 - accuracy: 0.5509 - val_loss: 1.1133 - val_accuracy: 0.5561
Epoch 6/50
31/31 [=====] - 19s 626ms/step - loss: 1.0707 - accuracy: 0.5688 - val_loss: 1.0658 - val_accuracy: 0.5831
Epoch 7/50
31/31 [=====] - 18s 655ms/step - loss: 1.0444 - accuracy: 0.5811 - val_loss: 1.0444 - val_accuracy: 0.5811
```

5. Loss and Accuracy check using plot

```
In [11]: #plot the loss
plt.plot(result.history['loss'], label='train loss')
plt.plot(result.history['val_loss'], label='val loss')
plt.legend()
plt.show()

# plot the accuracy
plt.plot(result.history['accuracy'], label='train acc')
plt.plot(result.history['val_accuracy'], label='val acc')
plt.legend()
plt.show()
```



6. Save the Model

```
In [12]: model.save('flower.h5')
```

Testing the Model

```
In [13]: training_set.class_indices
```

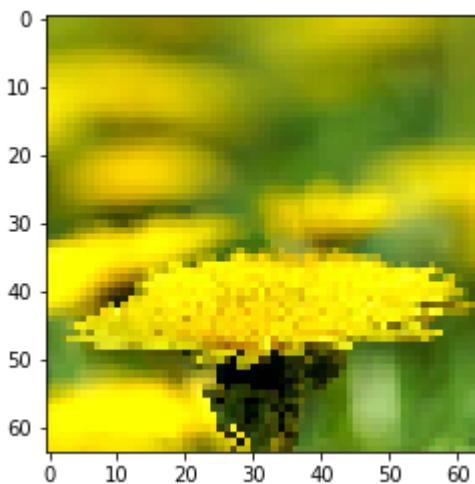
```
Out[13]: {'daisy': 0, 'dandelion': 1, 'rose': 2, 'sunflower': 3, 'tulip': 4}
```

```
In [26]: classes = ['Daisy', 'Dandelion', 'Rose', 'Sunflower', 'Tulip']
def testing(img):
    img = image.load_img(img, target_size=(64, 64))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    pred = np.argmax(model.predict(x))
    return print("Predicted class as:", classes[pred])

def img_show(img):
    img1 = image.load_img(img, target_size=(64, 64))
    plt.imshow(img1)
```


```
In [35]: #test1
img_show('flower1.jpg')
testing('flower1.jpg')
```

Predicted class as: Dandelion




```
In [31]: #test2
img_show('flower2.jpg')
testing('flower2.jpg')
```

Predicted class as: Rose




```
In [30]: #test3  
img_show('flower3.jpg')  
testing('flower3.jpg')
```

Predicted class as: Tulip



```
In [32]: #test4  
img_show('flower4.jpg')  
testing('flower4.jpg')
```

Predicted class as: Sunflower



```
In [33]: #test5  
img_show('flower5.jpg')  
testing('flower5.jpg')
```

Predicted class as: Daisy



Conclusion:

- The dataset has about 4317 images from 5 different classes.
- Each classes have more than 500 images for training the data.
- 30% of the data taken for validation.
- The accuracy of the model is around 80%.
- The validation accuracy is around 70%.
- The model is built with 2 layered convolutional network considering 1344 trainable parameters.
- Testing the model with unknown images gives 95% accuracy.