

Assignment 3 - Mythili K

Dataset

Download the dataset [here](#)

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Convolution2D, MaxPooling2D, Flatten, Dense
```

Data augmentation

```
In [2]: # Data augmentaion on training variable

train_datagen = ImageDataGenerator(rescale = 1./255 , zoom_range = 0.2 , horizontal_flip =
```

```
In [3]: # Data augmentation on testing varibale

test_datagen = ImageDataGenerator(rescale = 1./255)
```

```
In [4]: # Data augmentation on training data

xtrain = train_datagen.flow_from_directory('dataset/train/' ,
                                           target_size = (64,64) ,
                                           class_mode = 'categorical' ,
                                           batch_size=100)
```

Found 3019 images belonging to 5 classes.

```
In [5]: # Data augmentation on testing data

xtest = test_datagen.flow_from_directory('dataset/test/' ,
                                         target_size=(64,64) ,
                                         class_mode='categorical' ,
                                         batch_size=100)
```

Found 438 images belonging to 5 classes.

Model Building

```
In [6]: # Initializing the sequential model
model = Sequential()

# Convolutional layer
model.add(Convolution2D(32, (3,3), activation='relu', input_shape=(64,64,3)))
```

```

# Maxpooling layer
model.add(MaxPooling2D(pool_size=(2,2)))

# Flatten layer
model.add(Flatten())

# Hidden layer 1
model.add(Dense(64,activation='relu'))

# Hidden layer 2
model.add(Dense(32,activation='relu'))

# Output layer
model.add(Dense(5,activation='softmax')) # output

```

Compiling the model

```

In [7]: model.compile(optimizer='adam',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])

```

Training the model

```

In [8]: out=model.fit(xtrain,
                    epochs=50,
                    validation_data = xtest,
                    validation_steps = len(xtest))

```

```

Epoch 1/50
31/31 [=====] - 73s 2s/step - loss: 1.5628 - accuracy: 0.3640 - v
al_loss: 1.2511 - val_accuracy: 0.4749
Epoch 2/50
31/31 [=====] - 68s 2s/step - loss: 1.2220 - accuracy: 0.4929 - v
al_loss: 1.1798 - val_accuracy: 0.4772
Epoch 3/50
31/31 [=====] - 66s 2s/step - loss: 1.1637 - accuracy: 0.5293 - v
al_loss: 1.0959 - val_accuracy: 0.5616
Epoch 4/50
31/31 [=====] - 35s 1s/step - loss: 1.0881 - accuracy: 0.5671 - v
al_loss: 1.1107 - val_accuracy: 0.5320
Epoch 5/50
31/31 [=====] - 23s 727ms/step - loss: 1.0322 - accuracy: 0.6055
- val_loss: 1.0971 - val_accuracy: 0.5639
Epoch 6/50
31/31 [=====] - 30s 972ms/step - loss: 1.0166 - accuracy: 0.6022
- val_loss: 1.0813 - val_accuracy: 0.5639
Epoch 7/50
31/31 [=====] - 26s 843ms/step - loss: 0.9718 - accuracy: 0.6293
- val_loss: 1.1142 - val_accuracy: 0.5525
Epoch 8/50
31/31 [=====] - 53s 2s/step - loss: 0.9519 - accuracy: 0.6350 - v
al_loss: 1.1438 - val_accuracy: 0.5342
Epoch 9/50
31/31 [=====] - 41s 1s/step - loss: 0.9531 - accuracy: 0.6313 - v
al_loss: 1.0313 - val_accuracy: 0.6142
Epoch 10/50
31/31 [=====] - 27s 857ms/step - loss: 0.9133 - accuracy: 0.6535
- val_loss: 1.0819 - val_accuracy: 0.5868
Epoch 11/50
31/31 [=====] - 25s 803ms/step - loss: 0.8945 - accuracy: 0.6555
- val_loss: 1.0212 - val_accuracy: 0.6301

```

Epoch 12/50
31/31 [=====] - 21s 682ms/step - loss: 0.8574 - accuracy: 0.6698
- val_loss: 0.9979 - val_accuracy: 0.6164
Epoch 13/50
31/31 [=====] - 26s 831ms/step - loss: 0.8378 - accuracy: 0.6870
- val_loss: 0.9670 - val_accuracy: 0.6324
Epoch 14/50
31/31 [=====] - 27s 856ms/step - loss: 0.8524 - accuracy: 0.6767
- val_loss: 1.0705 - val_accuracy: 0.5890
Epoch 15/50
31/31 [=====] - 31s 1s/step - loss: 0.8255 - accuracy: 0.6843 - v
al_loss: 0.9381 - val_accuracy: 0.6507
Epoch 16/50
31/31 [=====] - 28s 901ms/step - loss: 0.7804 - accuracy: 0.7079
- val_loss: 1.0082 - val_accuracy: 0.6324
Epoch 17/50
31/31 [=====] - 37s 1s/step - loss: 0.7846 - accuracy: 0.7098 - v
al_loss: 0.8904 - val_accuracy: 0.6621
Epoch 18/50
31/31 [=====] - 59s 2s/step - loss: 0.7533 - accuracy: 0.7088 - v
al_loss: 0.9675 - val_accuracy: 0.6507
Epoch 19/50
31/31 [=====] - 48s 2s/step - loss: 0.7581 - accuracy: 0.7128 - v
al_loss: 1.0419 - val_accuracy: 0.6119
Epoch 20/50
31/31 [=====] - 60s 2s/step - loss: 0.7270 - accuracy: 0.7337 - v
al_loss: 0.9127 - val_accuracy: 0.6781
Epoch 21/50
31/31 [=====] - 28s 893ms/step - loss: 0.7094 - accuracy: 0.7314
- val_loss: 0.9735 - val_accuracy: 0.6256
Epoch 22/50
31/31 [=====] - 28s 899ms/step - loss: 0.6710 - accuracy: 0.7575
- val_loss: 1.0004 - val_accuracy: 0.6393
Epoch 23/50
31/31 [=====] - 57s 2s/step - loss: 0.6749 - accuracy: 0.7559 - v
al_loss: 0.9598 - val_accuracy: 0.6553
Epoch 24/50
31/31 [=====] - 65s 2s/step - loss: 0.6531 - accuracy: 0.7519 - v
al_loss: 0.9218 - val_accuracy: 0.6621
Epoch 25/50
31/31 [=====] - 54s 2s/step - loss: 0.6728 - accuracy: 0.7473 - v
al_loss: 0.9782 - val_accuracy: 0.6484
Epoch 26/50
31/31 [=====] - 58s 2s/step - loss: 0.6206 - accuracy: 0.7661 - v
al_loss: 0.9282 - val_accuracy: 0.6575
Epoch 27/50
31/31 [=====] - 29s 919ms/step - loss: 0.6245 - accuracy: 0.7632
- val_loss: 1.1047 - val_accuracy: 0.6416
Epoch 28/50
31/31 [=====] - 21s 687ms/step - loss: 0.6110 - accuracy: 0.7688
- val_loss: 0.9691 - val_accuracy: 0.6461
Epoch 29/50
31/31 [=====] - 29s 935ms/step - loss: 0.5886 - accuracy: 0.7781
- val_loss: 0.9532 - val_accuracy: 0.6553
Epoch 30/50
31/31 [=====] - 28s 891ms/step - loss: 0.5938 - accuracy: 0.7864
- val_loss: 0.9264 - val_accuracy: 0.6484
Epoch 31/50
31/31 [=====] - 22s 713ms/step - loss: 0.5659 - accuracy: 0.7917
- val_loss: 1.0002 - val_accuracy: 0.6370
Epoch 32/50
31/31 [=====] - 27s 865ms/step - loss: 0.5519 - accuracy: 0.7917
- val_loss: 0.9586 - val_accuracy: 0.6484
Epoch 33/50
31/31 [=====] - 23s 727ms/step - loss: 0.5437 - accuracy: 0.7993
- val_loss: 1.2153 - val_accuracy: 0.6119

```

Epoch 34/50
31/31 [=====] - 23s 740ms/step - loss: 0.5418 - accuracy: 0.7996
- val_loss: 1.0377 - val_accuracy: 0.6530
Epoch 35/50
31/31 [=====] - 25s 822ms/step - loss: 0.5082 - accuracy: 0.8168
- val_loss: 0.9776 - val_accuracy: 0.6530
Epoch 36/50
31/31 [=====] - 27s 854ms/step - loss: 0.5012 - accuracy: 0.8175
- val_loss: 1.1493 - val_accuracy: 0.6142
Epoch 37/50
31/31 [=====] - 28s 906ms/step - loss: 0.5106 - accuracy: 0.8198
- val_loss: 1.0325 - val_accuracy: 0.6484
Epoch 38/50
31/31 [=====] - 32s 1s/step - loss: 0.5288 - accuracy: 0.8046 - v
al_loss: 1.0109 - val_accuracy: 0.6667
Epoch 39/50
31/31 [=====] - 26s 854ms/step - loss: 0.5031 - accuracy: 0.8185
- val_loss: 1.0826 - val_accuracy: 0.6370
Epoch 40/50
31/31 [=====] - 25s 797ms/step - loss: 0.4808 - accuracy: 0.8331
- val_loss: 1.0646 - val_accuracy: 0.6553
Epoch 41/50
31/31 [=====] - 28s 911ms/step - loss: 0.4581 - accuracy: 0.8248
- val_loss: 1.1024 - val_accuracy: 0.6461
Epoch 42/50
31/31 [=====] - 25s 804ms/step - loss: 0.4672 - accuracy: 0.8258
- val_loss: 1.1276 - val_accuracy: 0.6461
Epoch 43/50
31/31 [=====] - 27s 872ms/step - loss: 0.4397 - accuracy: 0.8433
- val_loss: 1.1588 - val_accuracy: 0.6530
Epoch 44/50
31/31 [=====] - 24s 771ms/step - loss: 0.4815 - accuracy: 0.8205
- val_loss: 1.1843 - val_accuracy: 0.6484
Epoch 45/50
31/31 [=====] - 25s 792ms/step - loss: 0.4629 - accuracy: 0.8304
- val_loss: 1.2142 - val_accuracy: 0.6370
Epoch 46/50
31/31 [=====] - 26s 840ms/step - loss: 0.4174 - accuracy: 0.8476
- val_loss: 1.1455 - val_accuracy: 0.6461
Epoch 47/50
31/31 [=====] - 24s 763ms/step - loss: 0.4114 - accuracy: 0.8546
- val_loss: 1.0965 - val_accuracy: 0.6621
Epoch 48/50
31/31 [=====] - 31s 1000ms/step - loss: 0.4060 - accuracy: 0.8543
- val_loss: 1.1322 - val_accuracy: 0.6461
Epoch 49/50
31/31 [=====] - 26s 843ms/step - loss: 0.3729 - accuracy: 0.8662
- val_loss: 1.1479 - val_accuracy: 0.6667
Epoch 50/50
31/31 [=====] - 25s 813ms/step - loss: 0.3540 - accuracy: 0.8768
- val_loss: 1.0904 - val_accuracy: 0.6689

```

Plotting loss and accuracy

In [9]:

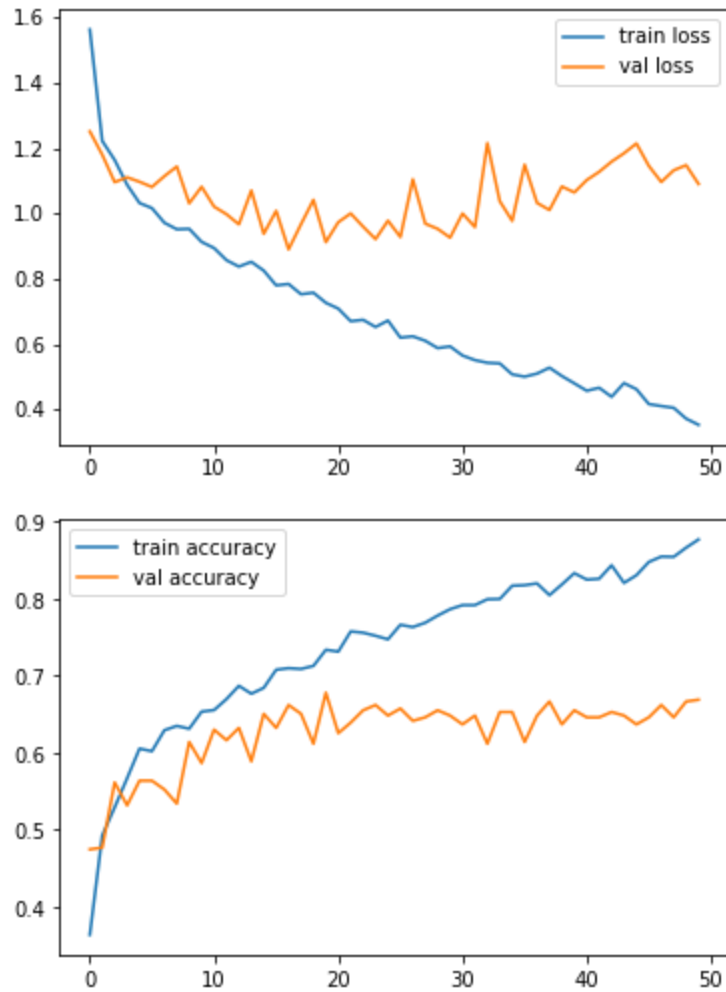
```

#plot the loss
plt.plot(out.history['loss'], label='train loss')
plt.plot(out.history['val_loss'], label='val loss')
plt.legend()
plt.show()

# plot the accuracy
plt.plot(out.history['accuracy'], label='train accuracy')
plt.plot(out.history['val_accuracy'], label='val accuracy')

```

```
plt.legend()  
plt.show()
```



```
In [10]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
flatten (Flatten)	(None, 30752)	0
dense (Dense)	(None, 64)	1968192
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 5)	165

```
=====  
Total params: 1,971,333  
Trainable params: 1,971,333  
Non-trainable params: 0
```

Save the model

```
In [11]:
```

```
model.save('flower.h5')
```

Testing the model


```
In [12]: # Creating list
flow = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']
def tester(img):
    img = image.load_img(img, target_size=(64, 64))
    # Converting images into array
    x = image.img_to_array(img)
    # Expanding the dimensions
    x = np.expand_dims(x, axis=0)
    # Predicting the higher probability index
    pred = np.argmax(model.predict(x))
    return print("Predicted class : ", flow[pred])

# Showing image
def show(img):
    img = image.load_img(img, target_size=(64, 64))
    plt.imshow(img)
```

Testing using flower images


```
In [13]: tester('flower1.jpg')
show('flower1.jpg')
```

Predicted class : tulip



```
In [14]: tester('flower2.jpg')
show('flower2.jpg')
```


Predicted class : sunflower



In [15]:

```
tester('flower5.jpg')  
show('flower5.jpg')
```

Predicted class : daisy



Inference

- The dataset comprises of five different classes of flowers with 4317 images
- The dataset is divided as 70% for training and 30% for testing and validation
- Model was built using Convolutional Neural network
- Accuracy : 85%
- Validation accuracy : 69%
- Testing accuracy : 90%